



[Return to Table of Contents](#)

Choose a Lesson

Choosing a Managed Database

Cloud SQL Basics

Importing Data

SQL Query Best Practices

Cloud SQL Basics

What is Cloud SQL?

- Direct lift and shift of traditional MySQL/PostgreSQL workloads with the maintenance stack managed for you

What is managed?

- OS installation/management
- Database installation/management
- Backups
- Scaling - disk space
- Availability:
 - Failover
 - Read replicas
- Monitoring
- Authorize network connections/proxy/use SSL

Limitations:

- Read replicas limited to the same region as the master:
 - Limited global availability
- Max disk size of 10 TB
- If > 10 TB is needed, or global availability in RDBMS, use **Spanner**

Scaling

High Availability

Database Backups

Software Patches

Database Installs

OS Patches

OS Installation

Server Maintenance

Physical Server

Power-Network-Cooling

Monitoring

import/export
CSV file
to instance



[Return to Table of Contents](#)

Choose a Lesson

Choosing a Managed Database

Cloud SQL Basics

Importing Data

SQL Query Best Practices

Importing Data

Importing data into Cloud SQL:

- Cloud Storage as a staging ground
- SQL dump/CSV file format

Export/Import process:

- 1 Export SQL dump/CSV file:
 - SQL dump file cannot contain **triggers, views, stored procedures**

need to remove them


- 2 Get dump/CSV file into Cloud Storage
- 3 Import from Cloud Storage into Cloud SQL instance

Best Practices:

- Use correct flags for dump file (--'flag_name'):
 - Databases, hex-blob, skip-triggers, set-gtid-purged=OFF, ignore-table
- Compress data to reduce costs:
 - Cloud SQL can import compressed .gz files
- Use InnoDB for Second Generation instances



SQL dump/CSV files



Cloud Storage



Cloud SQL

ground staging

load

2

3



[Return to Table of Contents](#)

Choose a Lesson

[Choosing a Managed Database](#)

[Cloud SQL Basics](#)

[Importing Data](#)

[SQL Query Best Practices](#)

SQL Query Best Practices

General SQL efficiency best practices:

- More, smaller tables better than fewer, large tables:
 - Normalization of tables
- Define your SELECT fields instead of using SELECT *:
 - SELECT * acts as a 'select all'
- When joining tables, use INNER JOIN instead of WHERE:
 - WHERE creates more variable combinations = more work

Questions:

SQL table grow over time
larger table → slow query



[Return to Table of Contents](#)

Choose a Lesson

[Cloud Datastore Overview](#)

[Data Organization](#)

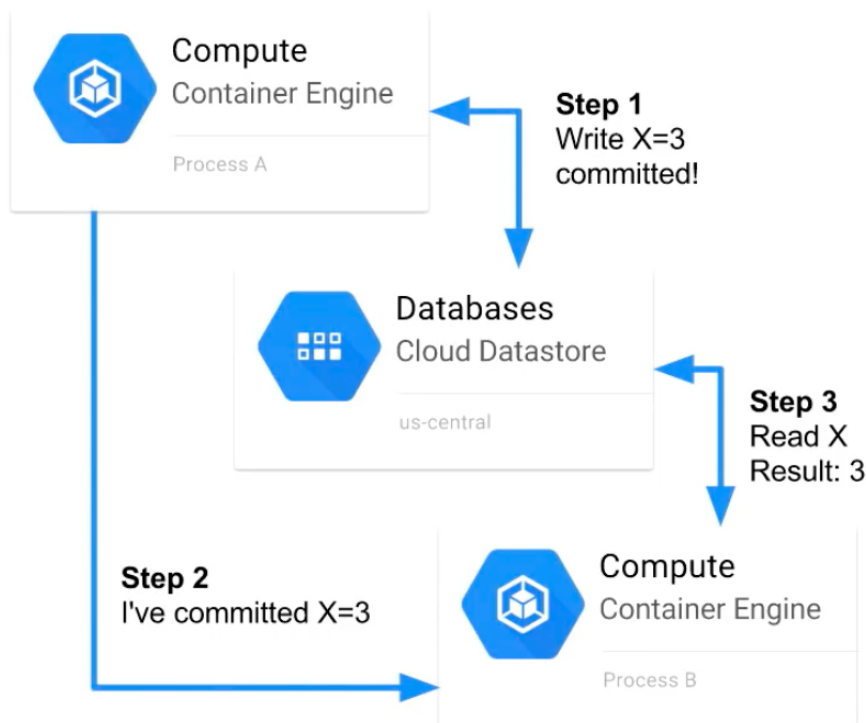
[Queries and Indexing](#)

[Data Consistency](#)

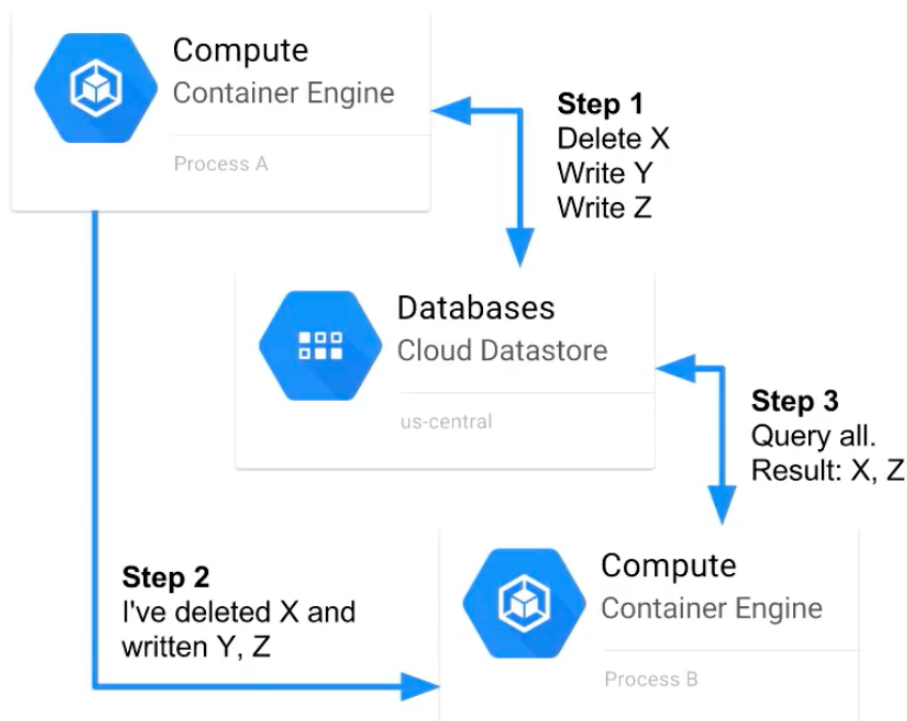
[Previous](#)

Data Consistency

Strong



Eventual





[Return to Table of Contents](#)

Choose a Lesson

[Cloud Datastore Overview](#)

[Data Organization](#)

[Queries and Indexing](#)

[Data Consistency](#)

(Cloud Datastore) **Data Consistency**

* Datastore is a NoOps Database

[Next](#)

What is data consistency in queries?

- "How up to date are these results?"
- "Does the order matter?"
- **Strongly consistent** = Parallel processes see changes in same order:
 - Query is **guaranteed** up to date but may take longer to complete
- **Eventually consistent** = Parallel process can see changes out of order, will eventually see accurate end state:
 - Faster query, but may *sometimes* **return stale results**
- Performance vs. accuracy
- Ancestor query/key-value operations = strong
- Global queries/projections = eventual

Use cases:

- Strong - financial transaction:
 - Make deposit -- check balance
- Eventual - census population:
 - Order not as important, as long as you get eventual result



[Return to Table of Contents](#)

Choose a Lesson

[Cloud Datastore Overview](#)

[Data Organization](#)

[Queries and Indexing](#)

[Data Consistency](#)

Queries and Indexing

[Previous](#)

Danger - Exploding Indexes!

- Default - create an entry for every possible combination of property values
- Results in higher storage and degraded performance
- Solutions:
 - Use a custom index.yaml file to narrow index scope
 - Do not index properties that don't need indexing

[Return to Table of Contents](#)

Choose a Lesson

[Cloud Datastore Overview](#)[Data Organization](#)[Queries and Indexing](#)[Data Consistency](#)[Next](#)

Queries and Indexing

Query:

- Retrieve an entity from Datastore that meets a set of conditions
- Query includes:
 - Entity kind *(table)*
 - Filters *← "where" in SQL*
 - Sort order
- Query methods:
 - 1 Programmatic
 - 2 Web console
 - 3 Google Query Language (GQL)

Indexing:

- Queries gets results from indexes:
 - Contain entity ^{row} keys specified by index properties
 - Updated to reflect changes
 - Correct query results available with no additional computation needed

Index types:

- Built-in - default option: *column*
 - Allows single property queries *at one time (if 2 queries) (cause error)*
- Composite - specified with an index configuration file (index.yaml):
 - gcloud datastore create-indexes index.yaml

```
indexes:
- kind: Task
  properties:
  - name: tags
  - name: created
- kind: Task
  properties:
  - name: collaborators
  - name: created
```

example



[Return to Table of Contents](#)

Data Organization

[Previous](#)

Choose a Lesson

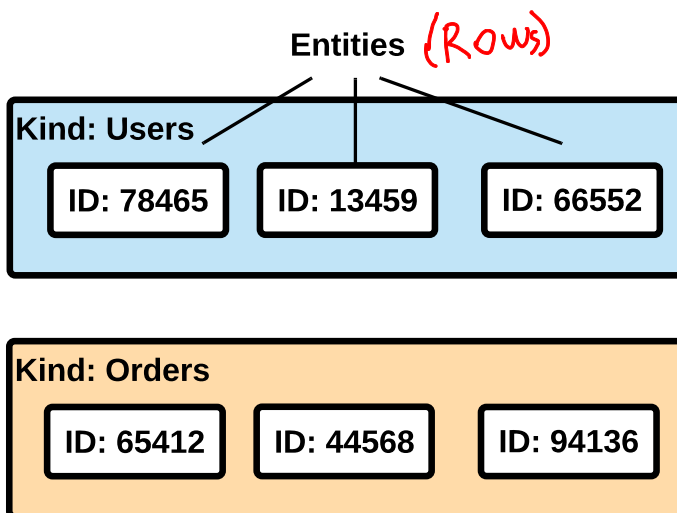
[Cloud Datastore Overview](#)

[Data Organization](#)

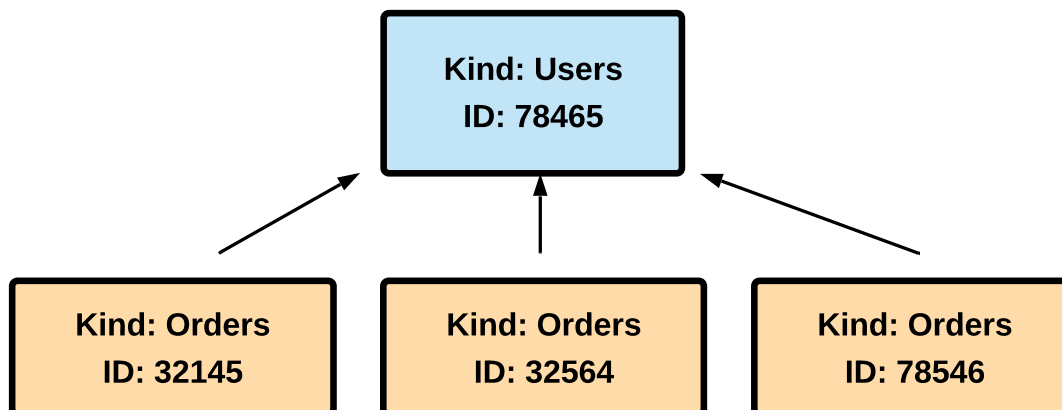
[Queries and Indexing](#)

[Data Consistency](#)

Simple Collections of Entities



Hierarchies (Entity Groups)



Return to Table of Contents

Choose a Lesson

Cloud Datastore Overview

Data Organization

Queries and Indexing

Data Consistency

Data Organization

Next

Short version:

- Entities grouped by kind (category)
- Entities can be hierarchical (nested)
- Each entity has one or more properties
- Properties have a value assigned

key: value store

| Concept | Relational Database | Datastore |
|-------------------------------|---------------------|-----------------|
| Category of object | Table | Kind (category) |
| Single Object | Row | Entity |
| Individual data for an object | Column | Property |
| Unique ID for an object | Primary key | Key |