

[Return to Table of Contents](#)

Choose a Lesson

[Streaming Data Challenges](#)[Cloud Pub/Sub Overview](#)[Pub/Sub Hands On](#)[Connecting Kafka to GCP](#)[Monitoring Subscriber Health](#)

Cloud Pub/Sub Overview

[Next](#)

What is Cloud Pub/Sub?

- Global-scale messaging buffer/coupler
- **NoOps**, global availability, auto-scaling
- **Decouples** senders and receivers
- Streaming data ingest:
 - Also connects other data pipeline services
- **Equivalent to Apache Kafka (open source)**
- Guaranteed at-least-once delivery

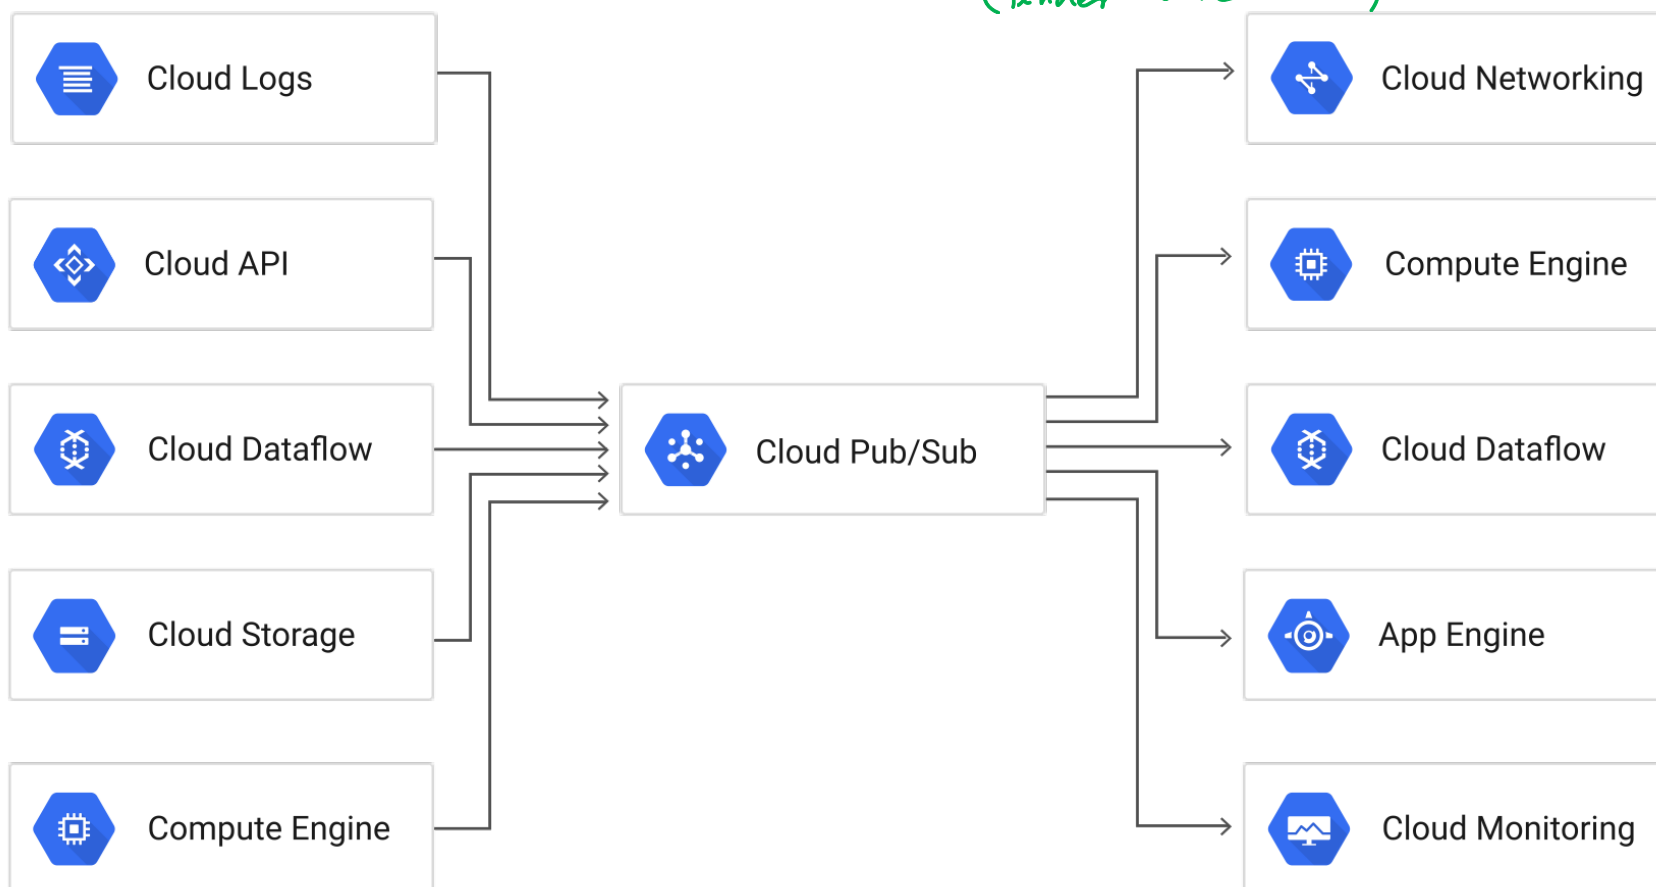
No ops } *No second figure*
No cluster to provision

exam question

Terminology ↘

Asynchronous messaging - many to many (or any other combination)
(Sender to receiver)

*1 to 1
many to many
7 to 8, 6 to 2, doesn't matter*



[Return to Table of Contents](#)

Cloud Pub/Sub Overview

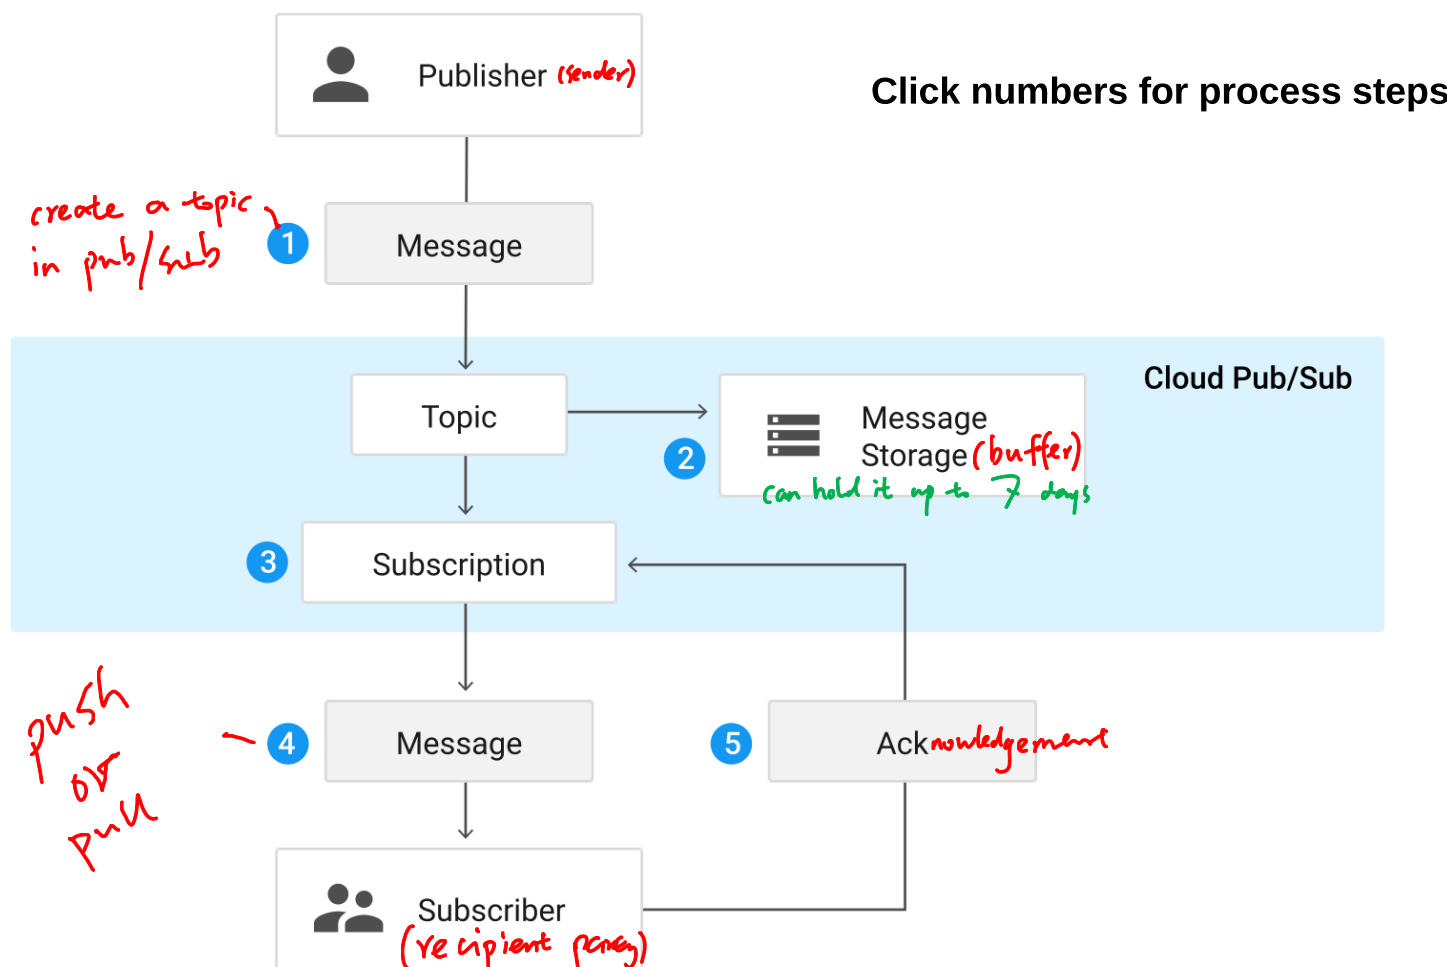
[Previous](#)[Next](#)

Choose a Lesson

[Streaming Data Challenges](#)[Cloud Pub/Sub Overview](#)[Pub/Sub Hands On](#)[Connecting Kafka to GCP](#)[Monitoring Subscriber Health](#)

How It Works: Terminology

- **Topics, Messages, Publishers, Subscribers, Message Store**



[Return to Table of Contents](#)

Cloud Pub/Sub Overview

[Previous](#)[Next](#)

Choose a Lesson

[Streaming Data Challenges](#)[Cloud Pub/Sub Overview](#)[Pub/Sub Hands On](#)[Connecting Kafka to GCP](#)[Monitoring Subscriber Health](#)

Push and Pull *(maybe in exam)*

- Pub/Sub can either **push messages to subscribers**, or **subscribers can pull messages from Pub/Sub**.
- **Push = lower latency, more real-time.**
- Push subscribers **must be Webhook** endpoints that accept POST over HTTPS.
- **Pull** is ideal for large volume of messages, and uses batch delivery
pull is the default option

IAM *(exam)*

- Allows for controlling access at project, topic, or subscription level
- Admin, Editor, Publisher, Subscriber
- **Service accounts** are best practice *(especially for built-ins in an application)*

Pricing

- Data volume used per month (per GB)

Monthly data	Price Per GB
First 10 GB	\$0.00
Next 50 TB	\$0.06
Next 100 TB	\$0.05
Beyond 150 TB	\$0.04

Out of order messaging

- Messages may **arrive from** multiple sources **out of order**.
- Pub/Sub **does not care** about message ordering.
- **Dataflow** is where out of order messages are **processed/resolved**.
- It's possible to add message attributes to help with ordering.

*(Pub/sub doesn't)**↑
exam*



[Return to Table of Contents](#)

Choose a Lesson

[Streaming Data Challenges](#)

[Cloud Pub/Sub Overview](#)

[Pub/Sub Hands On](#)

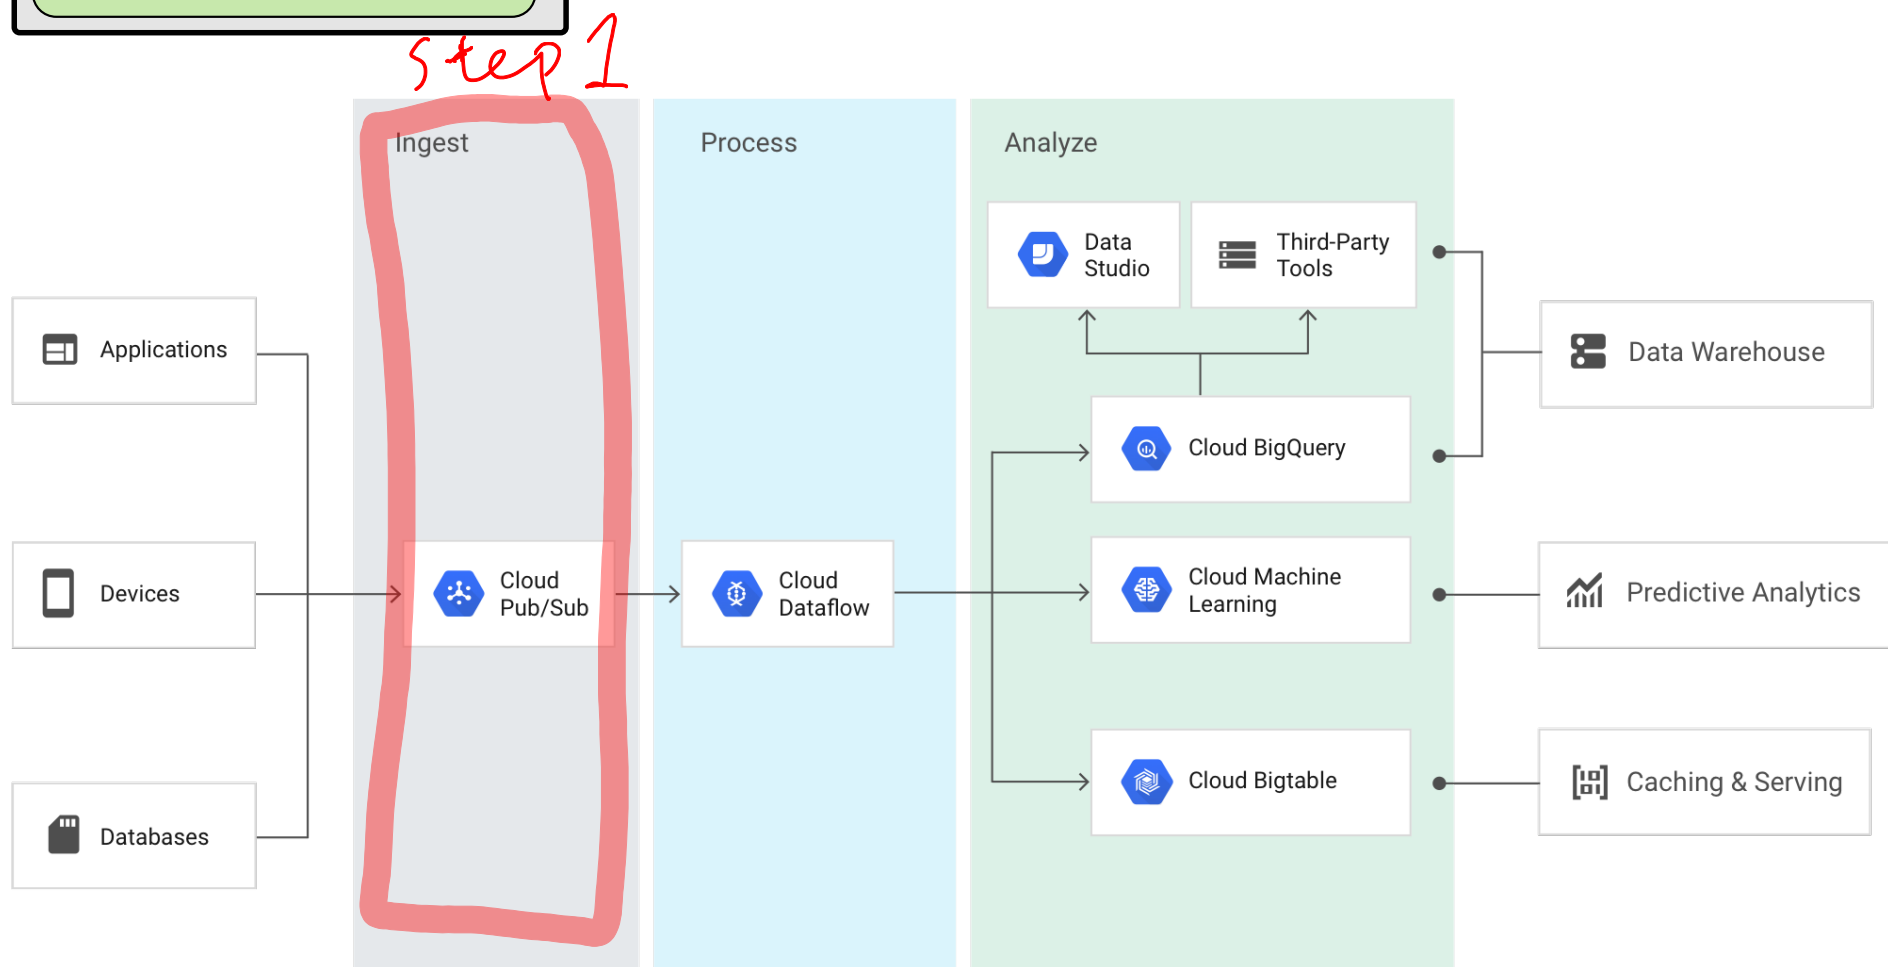
[Connecting Kafka to GCP](#)

[Monitoring Subscriber Health](#)

[Previous](#)

Cloud Pub/Sub Overview

Big Picture: Data Lifecycle for Streaming Data Ingest



[Return to Table of Contents](#)

Choose a Lesson

[Streaming Data Challenges](#)[Cloud Pub/Sub Overview](#)[Pub/Sub Hands On](#)[Connecting Kafka to GCP](#)[Monitoring Subscriber Health](#)

Pub/Sub Hands On

[Next](#)

The Steps

- Create a topic
- Create a subscription
- Publish messages
- Retrieve messages

Simple topic/subscription/publish via `gcloud`

Create a topic called *my-topic*:

- `gcloud pubsub topics create my-topic`

Create subscription to topic *my-topic*:

- `gcloud pubsub subscriptions create --topic my-topic mySub1`

Publish a message to your topic:

- `gcloud pubsub topics publish my-topic --message "hello"`

Retrieve message with your subscription, acknowledge receipt, and remove message from queue:

- `gcloud pubsub subscriptions pull --auto-ack mySub1`

Cancel subscription:

- `gcloud pubsub subscriptions delete mySub1`

[Return to Table of Contents](#)

Choose a Lesson

[Streaming Data Challenges](#)[Cloud Pub/Sub Overview](#)[Pub/Sub Hands On](#)[Connecting Kafka to GCP](#)[Monitoring Subscriber Health](#)

Pub/Sub Hands On

[Previous](#)

Traffic Data Exercise

- Clone GitHub
- Copy data points
- Simulate traffic data
- Pull messages

Clone GitHub data to Cloud Shell (or other SDK environment), and browse to publish folder:

```
cd ~
git clone https://github.com/linuxacademy/googledataengineer
cd ~/googledataengineer/courses/streaming/publish
```

Create a topic called *sandiego*:

```
gcloud pubsub topics create sandiego
```

Create subscription to topic *sandiego*:

```
gcloud pubsub subscriptions create --topic sandiego mySub1
```

Run script to download sensor data:

```
./download_data.sh
```

May need to authenticate shell to ensure we have the right permissions:

```
gcloud auth application-default login
```

View script info:

```
vim ./send_sensor_data.py or use viewer of your choice
```

Run python script to simulate one hour of data per minute:

```
./send_sensor_data.py --speedFactor=60 \
--project=YOUR-PROJECT-ID
```

If you receive error: *google.cloud.pubsub can not be found* or an *ImportError: No module named iterator*, run this **pip** command to install components, then try again:

```
sudo pip install -U google-cloud-pubsub
```

Open new Cloud Shell tab (using + symbol)

Pull message using subscription *mySub1*:

```
gcloud pubsub subscriptions pull --auto-ack mySub1
```

Create a new subscription and pull messages with it:

```
gcloud pubsub subscriptions create --topic sandiego mySub2
gcloud pubsub subscriptions pull --auto-ack mySub2
```

[Return to Table of Contents](#)

Choose a Lesson

[Streaming Data Challenges](#)[Cloud Pub/Sub Overview](#)[Pub/Sub Hands On](#)[Connecting Kafka to GCP](#)[Monitoring Subscriber Health](#)

Connecting Kafka to GCP

[Next](#)

Does Pub/Sub Replace Kafka?

- Not always
- Hybrid workloads:
 - Interact with existing tools and frameworks
 - Don't need global/scaling capabilities with pub/sub
- Can use *both*: **Kafka for on-premises** and **pub/sub for GCP** in same data pipeline

How do we connect Kafka to GCP?

Overview on **Connectors**:

- Open-source **plugins** that connect Kafka to GCP
- Kafka Connect: One optional "connector service"
- Exist to connect Kafka **directly** to **pub/sub**, **Dataflow**, and **BigQuery** (among others)

Additional Terms

- **Source connector**: An upstream connector:
 - Streams *from* something **to** Kafka
- **Sink connector**: A downstream connector:
 - Streams *from* Kafka **to** something else

exam related terminology



Source
Connector



kafka

Sink
Connector



Cloud Pub/Sub

On-Premises
Database



[Return to Table of Contents](#)

Choose a Lesson

Streaming Data Challenges

Cloud Pub/Sub Overview

Pub/Sub Hands On

Connecting Kafka to GCP

Monitoring Subscriber Health

Monitoring Subscriber Health


Trouble shooting

In a Perfect World...

- Subscribers and Publishers work in perfect harmony:
 - Example:
 - 1 million messages/second published
 - 1 million messages/second successfully pulled/pushed
 - Result: No backlog in Pub/Sub queue
- But we don't live in a perfect world...
 - Subscriber cannot keep up with publish rate
 - Result: Backlog in Pub/Sub queue

Troubleshooting Subscriber Health (Backlog)

- Create alerts for (x) backlog threshold
- Subscriber not able to keep up:
 - Under-provisioned
 - Code not optimized
- Not acknowledging message receipt:
 - Pub/Sub doesn't know it's delivered, and keeps trying
 - Subscriber code not properly acknowledging pulled messages
- Check publishers for excessive re-transmits

*4 trouble
shootings**two reasons**acknowledging Requests
= Pull Message Operations**(publishers)*

Sensors




1 million/sec



Cloud Pub/Sub



10,000/sec

Cloud
Dataflow*(subscribers)*